

Lab 7: Case study with Keras

authors: Arias Rodrigo, Burca Horia

date: 22/11/2018

In order to run keras in Marenosturm, we first need to load some modules. The script `load-modules` contains the needed modules, depending on the cluster:

```
#!/bin/bash

if [ $(hostname) == 'p9login1' ]; then
    module load cudnn/7.1.3 atlas/3.10.3 scalapack/2.0.2 fftw/3.3.7 \
        szip/2.1.1 opencv/3.4.1 python/3.6.5_ML
else
    module load gcc/7.2.0 impi/2018.1 mkl/2018.1 python/3.6.4_ML
fi
```

Load the required modules by running:

```
$ . load-modules
```

To run, use python3:

```
$ python3 xx.py
```

Loading the dataset

The `mnist` dataset was manually placed in the Marenosturm cluster, and then loaded in keras. The file `train.py` trains the model and saves the state to disk.

First we load the dataset:

```
#!/usr/bin/env python
# coding: utf-8

from keras.datasets import mnist
from keras.layers import *
from keras.models import *
from keras.utils import to_categorical

# Save the files manually in ~/.keras/datasets/ as MN cannot download from
# Internet
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# print(x_test.shape)
# (10000, 28, 28)

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)
```

```
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

To build the model, we use the Sequential class:

```
# Build the model

model = Sequential()
model.add(Conv2D(32, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(10, activation='softmax'))

model.summary()
```

And then we train the model:

```
batch_size = 50
num_classes = 10
epochs=50

model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1)

test_loss, test_acc = model.evaluate(x_test, y_test)

print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

Finally, the model is saved to disk, with the weights:

```
#Serialize model to JSON
model_json = model.to_json()
with open("model.json", "w") as f:
    f.write(model_json)
# serialize weights to HDF5
model.save_weights("model.h5")
print("Saved model to disk")
```

Prediction and evaluation

After training the model, we load again the dataset:

```
#!/usr/bin/env python
# coding: utf-8

import keras
```

```

import matplotlib
matplotlib.use('Agg')
import matplotlib.pyplot as plt
import numpy as np
from numpy import linalg
from keras.datasets import mnist
from keras.layers import *
from keras.models import *
from keras.optimizers import sgd
from keras.utils import to_categorical

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# print(x_test.shape)
# (10000, 28, 28)

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)

y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)

```

Then we load the trained model:

```

# load json and create model
json_file = open('model.json', 'r')
model_json = json_file.read()
json_file.close()
model = model_from_json(model_json)
# load weights into new model
model.load_weights("model.h5")
print("Loaded model from disk")

```

And evaluate the accuracy:

```

# evaluate loaded model on test data
model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

test_loss, test_acc = model.evaluate(x_test, y_test)

print('Test loss:', test_loss)
print('Test accuracy:', test_acc)

```

Which surpasses the provided model:

```

Test loss: 0.04247535159343388
Test accuracy: 0.9859

```


And an example is tested

```
# Predictions

i = 11
predictions = model.predict(x_test)
tag = np.argmax(predictions[i])

test_sel = x_test[i]
img = test_sel.reshape(28, 28)
plt.imshow(img)
plt.title("Example classified as {}".format(tag))
plt.savefig("test.png")
plt.close()
```

Which is correctly classified:



Finally, the confusion matrix is plotted as well:

```
# Look at confusion matrix
# Note, this code is taken straight from the SKLEARN website, an nice way of
# viewing confusion matrix.

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
```

```

plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, cm[i, j],
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('Actual class')
plt.xlabel('Predicted class')

from collections import Counter
from sklearn.metrics import confusion_matrix
import itertools

# Predict the values from the validation dataset
Y_pred = model.predict(x_test)
# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred, axis = 1)
# Convert validation observations to one hot vectors
Y_true = np.argmax(y_test, axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)
# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes = range(10))
plt.savefig("plot.png")

```

