

Lab 7: Case study with Keras

authors: Arias Rodrigo, Burca Horia
date: 22/11/2018

In order to run keras in Marenosturm, we first need to load some modules. The script load-modules contains the needed modules, depending on the cluster:

```
#!/bin/bash

if [ $(hostname) == 'p9login1' ]; then
    module load cudnn/7.1.3 atlas/3.10.3 scalapack/2.0.2 fftw/3.3.7 \
        szip/2.1.1 opencv/3.4.1 python/3.6.5_ML
else
    module load gcc/7.2.0 impi/2018.1 mkl/2018.1 python/3.6.4_ML
fi
```

Load the required modules by running:

```
$ . load-modules
```

To run, use python3:

```
$ python3 xx.py
```

Loading the dataset

The mnist dataset was manually placed in the Marenosturm cluster, and then loaded in keras. The file train.py trains the model and saves the state to disk.

First we load the dataset:

```
#!/usr/bin/env python
# coding: utf-8

from keras.datasets import mnist
from keras.layers import *
from keras.models import *
from keras.utils import to_categorical

# Save the files manually in ~/.keras/datasets/ as MN cannot download from
# Internet
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# print(x_test.shape)
```

```
# (10000, 28, 28)

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)

y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

To build the model, we use the Sequential class:

```
# Build the model

model = Sequential()
model.add(Conv2D(32, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(10, activation='softmax'))

model.summary()
```

And then we train the model:

```
batch_size = 50
num_classes = 10
epochs=50

model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1)

test_loss, test_acc = model.evaluate(x_test, y_test)

print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

Finally, the model is saved to disk, with the weights:

```
#Serialize model to JSON
model_json = model.to_json()
with open("model.json", "w") as f:
    f.write(model_json)
# serialize weights to HDF5
model.save_weights("model.h5")
print("Saved model to disk")
```

Prediction and evaluation

After training the model, we load again the dataset:

```
#!/usr/bin/env python
# coding: utf-8

import keras
import matplotlib
matplotlib.use('Agg')
import matplotlib.pyplot as plt
import numpy as np
from numpy import linalg
from keras.datasets import mnist
from keras.layers import *
from keras.models import *
from keras.optimizers import sgd
from keras.utils import to_categorical

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# print(x_test.shape)
# (10000, 28, 28)

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)

y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

Then we load the trained model:

```
# load json and create model
json_file = open('model.json', 'r')
model_json = json_file.read()
json_file.close()
model = model_from_json(model_json)
# load weights into new model
model.load_weights("model.h5")
print("Loaded model from disk")
```

And evaluate the accuracy:

```
# evaluate loaded model on test data
model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

test_loss, test_acc = model.evaluate(x_test, y_test)
```

```
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

Which surpasses the provided model:

```
Test loss: 0.04247535159343388
Test accuracy: 0.9859
```

And an example is tested

```
# Predictions

i = 11
predictions = model.predict(x_test)
tag = np.argmax(predictions[i])

test_sel = x_test[i]
img = test_sel.reshape(28, 28)
plt.imshow(img)
plt.title("Example classified as {}".format(tag))
plt.savefig("test.png")
plt.close()
```

Which is correctly classified:



Finally, the confusion matrix is plotted as well:

```
# Look at confusion matrix
# Note, this code is taken straight from the SKLEARN website, an nice way of
# viewing confusion matrix.

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
```

```
thresh = cm.max() / 2.  
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):  
    plt.text(j, i, cm[i, j],  
             horizontalalignment="center",  
             color="white" if cm[i, j] > thresh else "black")
```

```
plt.tight_layout()  
plt.ylabel('Actual class')  
plt.xlabel('Predicted class')
```

```
from collections import Counter  
from sklearn.metrics import confusion_matrix  
import itertools
```

```
# Predict the values from the validation dataset  
Y_pred = model.predict(x_test)  
# Convert predictions classes to one hot vectors  
Y_pred_classes = np.argmax(Y_pred, axis = 1)  
# Convert validation observations to one hot vectors  
Y_true = np.argmax(y_test, axis = 1)  
# compute the confusion matrix  
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)  
# plot the confusion matrix  
plot_confusion_matrix(confusion_mtx, classes = range(10))  
plt.savefig("plot.png")
```

